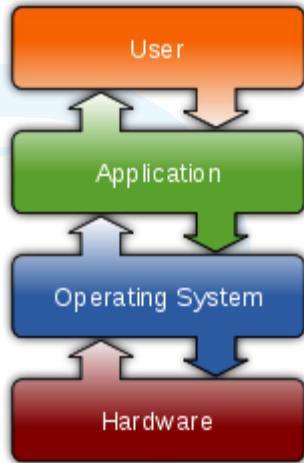


ENCE360

Operating Systems



Deadlocks

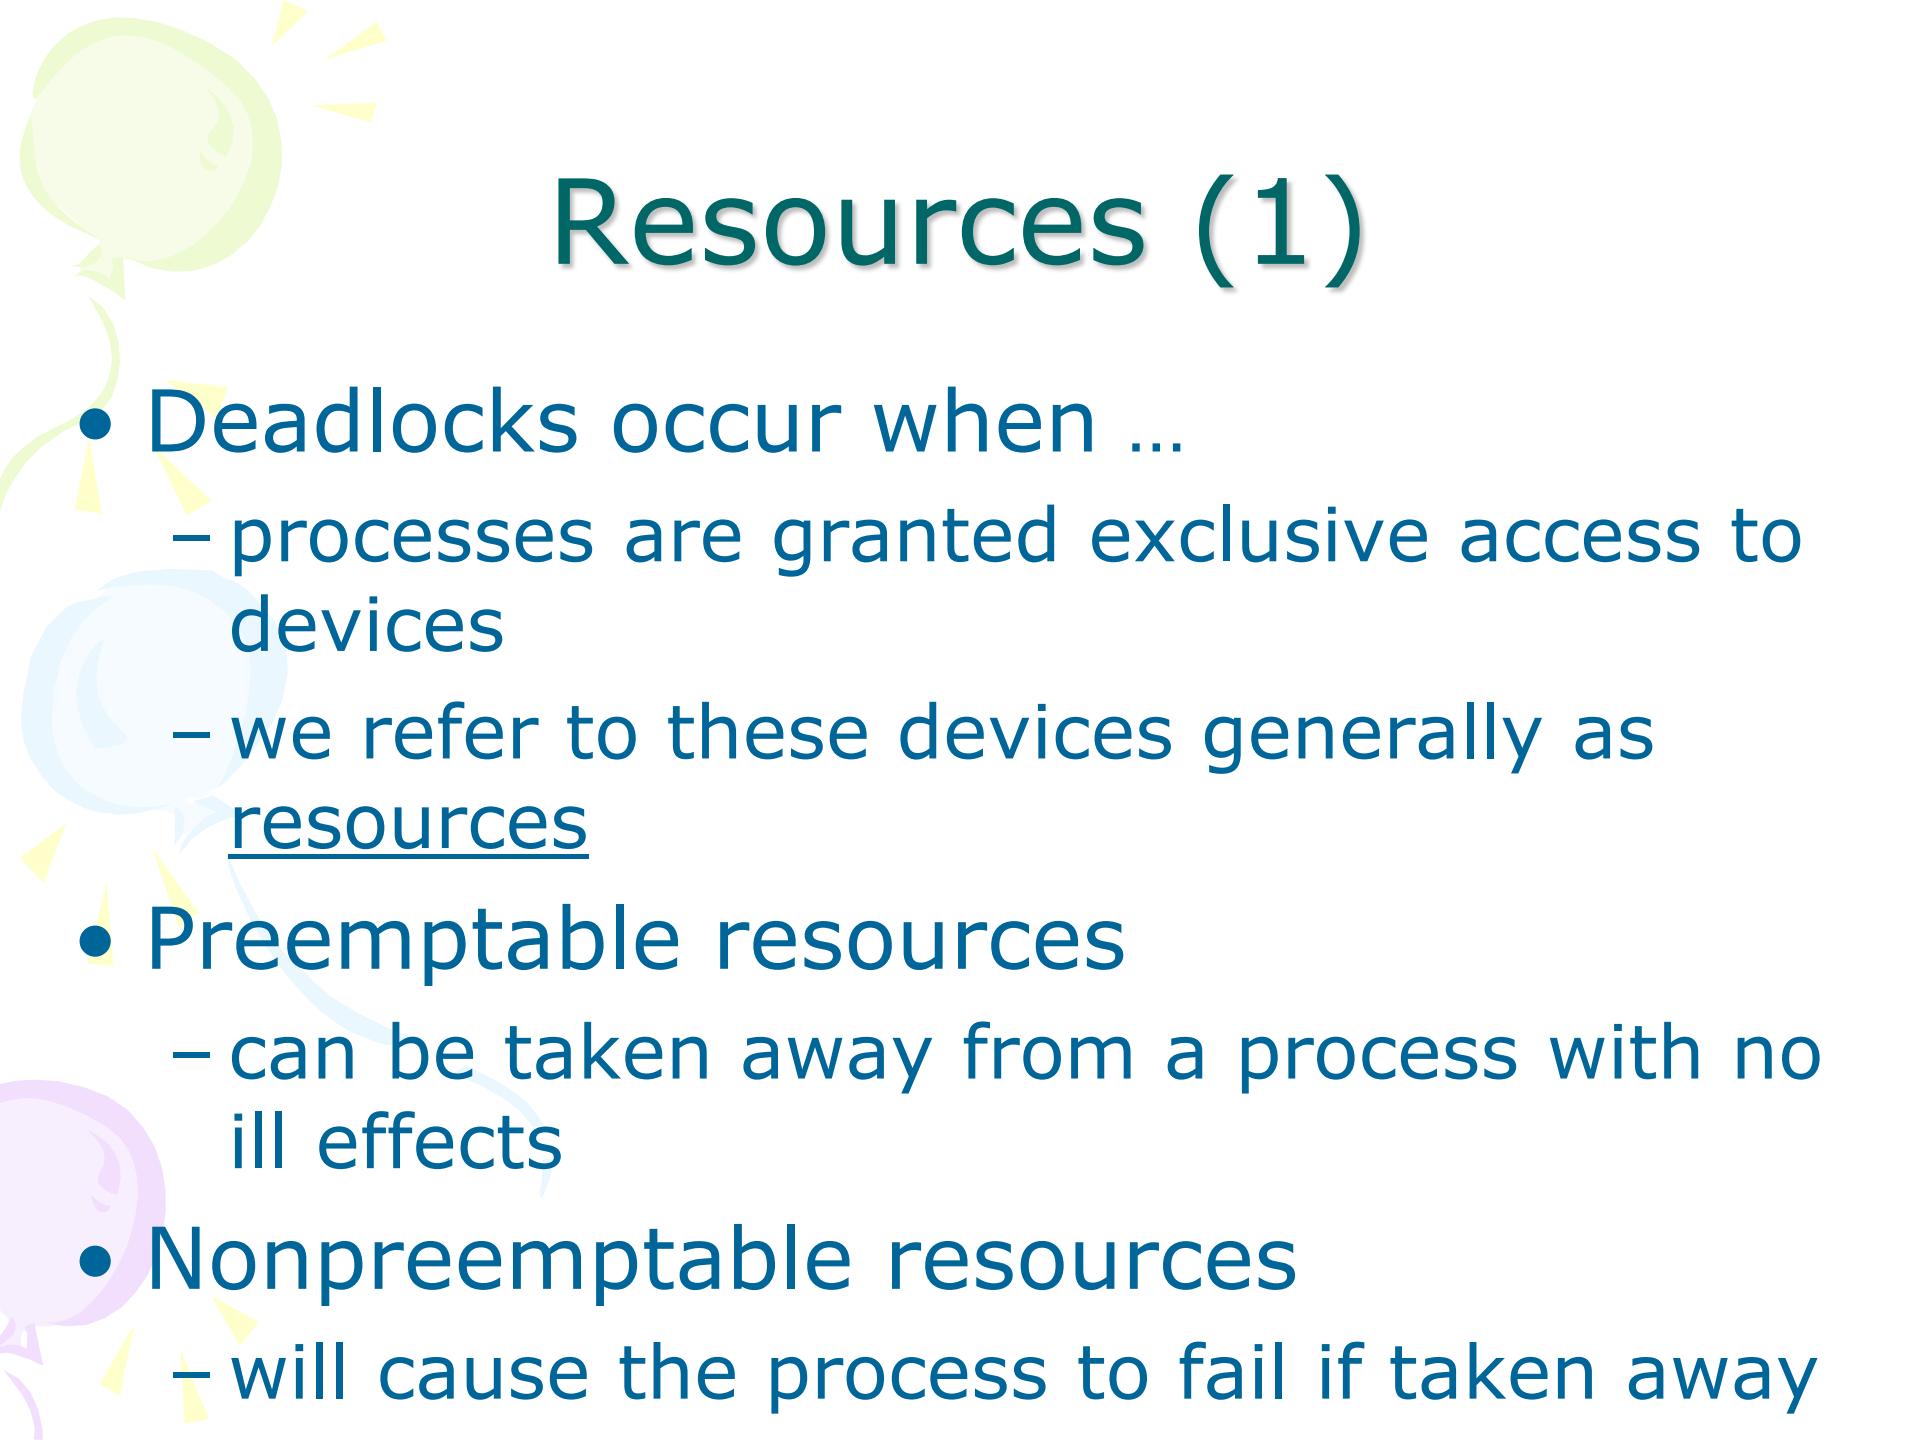
- Resource
- Introduction to deadlocks
- The ostrich algorithm
- Deadlock detection and recovery
- Deadlock avoidance
- Deadlock prevention

MOS Ch 3



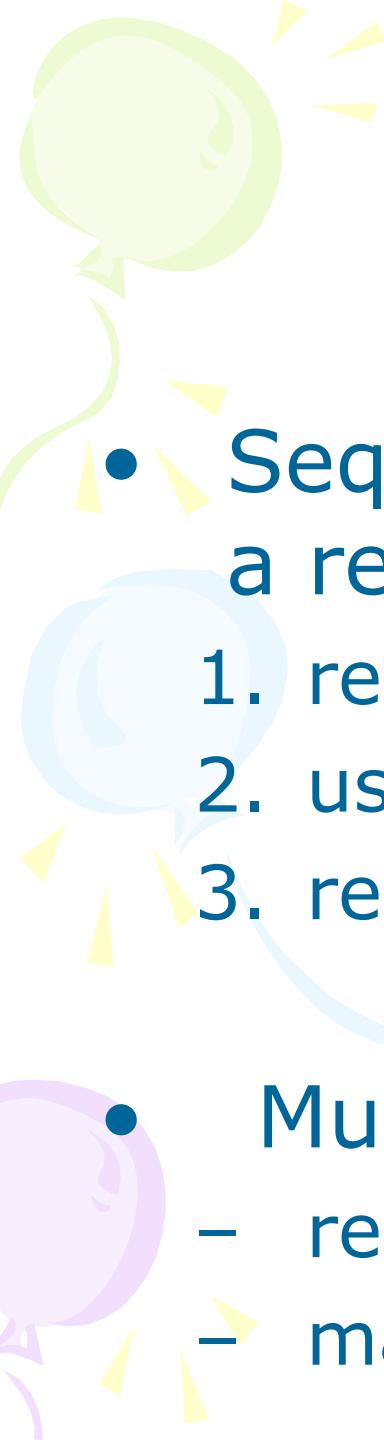
Resources

- Examples of computer resources
 - printers
 - tape drives
 - tables
- Processes need access to resources in reasonable order
- Suppose a process holds resource A and requests resource B
 - at same time another process holds B and requests A
 - both are blocked and remain so



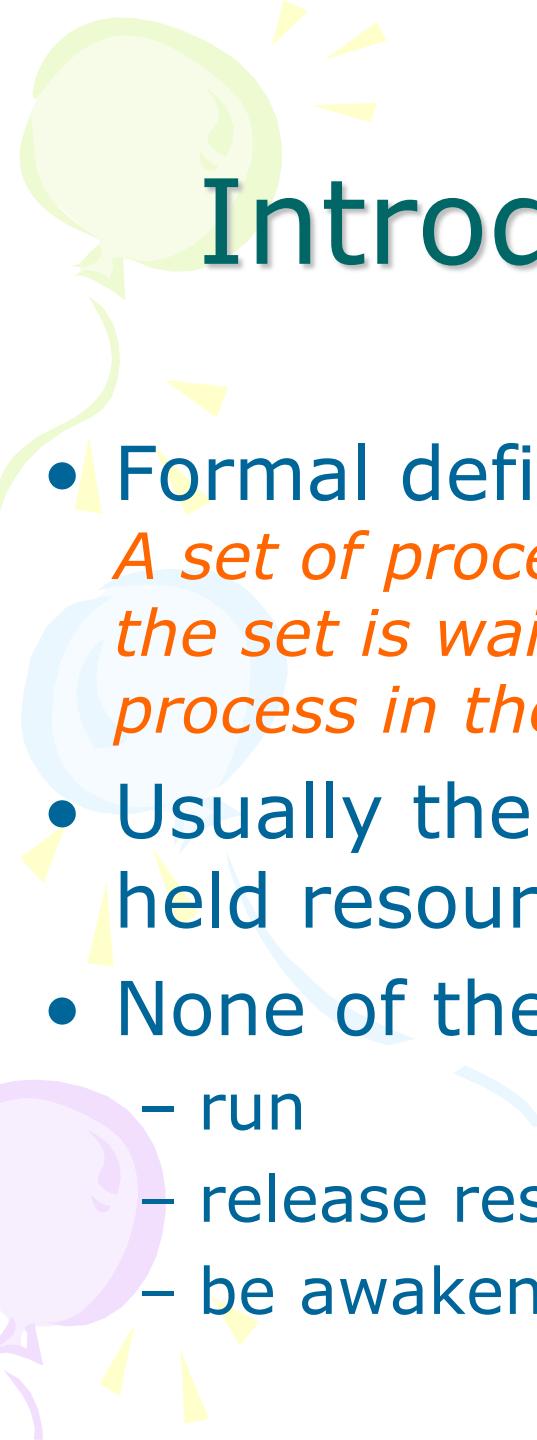
Resources (1)

- Deadlocks occur when ...
 - processes are granted exclusive access to devices
 - we refer to these devices generally as resources
- Preemptable resources
 - can be taken away from a process with no ill effects
- Nonpreemptable resources
 - will cause the process to fail if taken away



Resources (2)

- Sequence of events required to use a resource
 1. request the resource
 2. use the resource
 3. release the resource
- Must wait if request is denied
 - requesting process may be blocked
 - may fail with error code



Introduction to Deadlocks

- Formal definition :

A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause

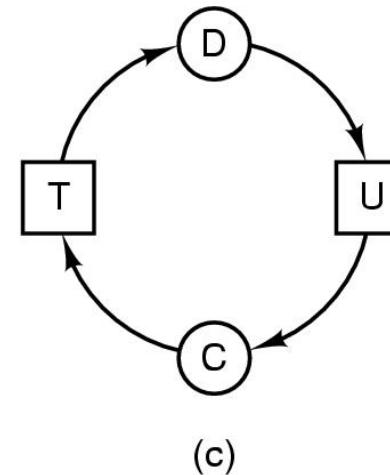
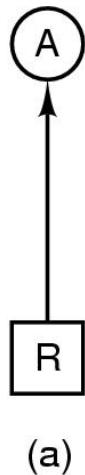
- Usually the event is release of a currently held resource
- None of the processes can ...
 - run
 - release resources
 - be awakened

Four Conditions for Deadlock

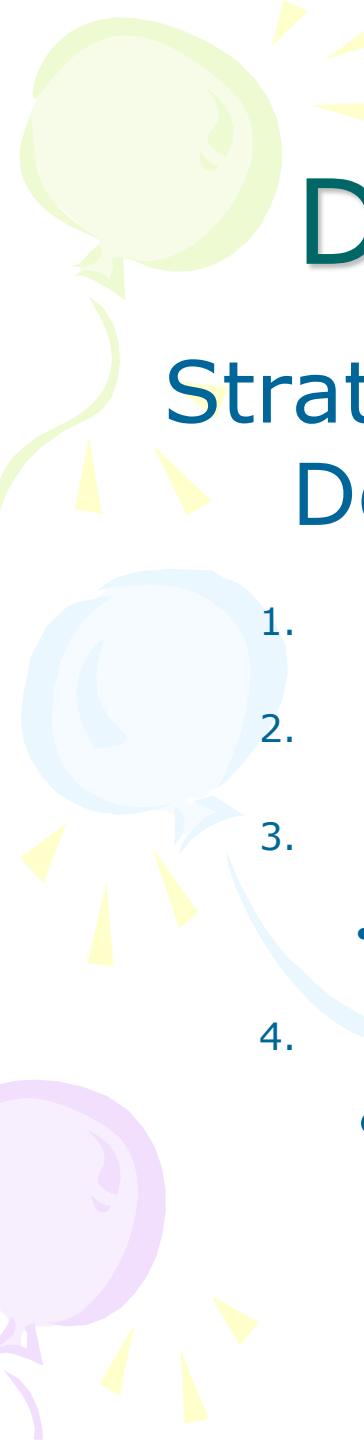
1. Mutual exclusion condition
 - each resource assigned to 1 process or is available
2. Hold and wait condition
 - process holding resources can request additional
3. No preemption condition
 - previously granted resources cannot forcibly taken away
4. Circular wait condition
 - must be a circular chain of 2 or more processes
 - each is waiting for resource held by next member of the chain

Deadlock Modeling (1)

- Modeled with directed graphs



- resource R assigned to process A
- process B is requesting/waiting for resource S
- process C and D are in deadlock over resources T and U



Deadlock Modeling (2)

Strategies for dealing with Deadlocks

1. just ignore the problem altogether
2. detection and recovery
3. dynamic avoidance
 - careful resource allocation
4. prevention
 - negating one of the four necessary conditions

Deadlock Modeling (3)

A

Request R
Request S
Release R
Release S

(a)

B

Request S
Request T
Release S
Release T

(b)

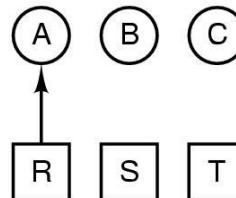
C

Request T
Request R
Release T
Release R

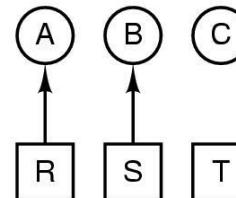
(c)

1. A requests R
 2. B requests S
 3. C requests T
 4. A requests S
 5. B requests T
 6. C requests R
- deadlock

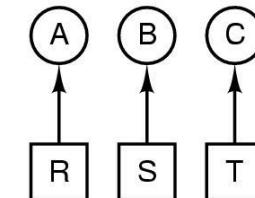
(d)



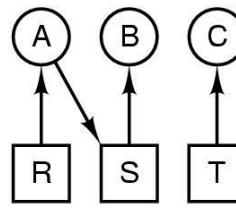
(e)



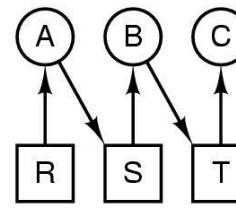
(f)



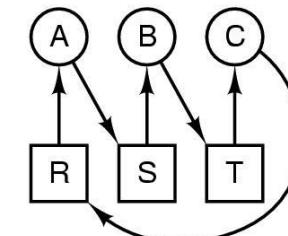
(g)



(h)



(i)



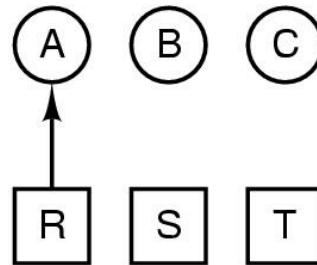
(j)

How deadlock occurs

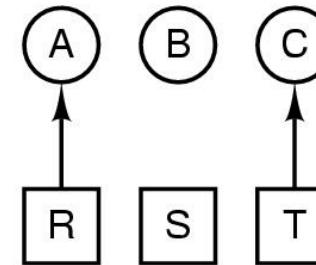
Deadlock Modeling (4)

1. A requests R
2. C requests T
3. A requests S
4. C requests R
5. A releases R
6. A releases S
- no deadlock

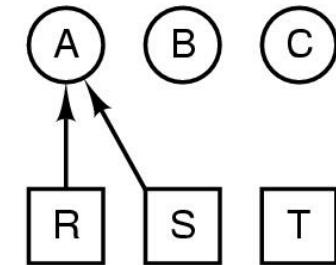
(k)



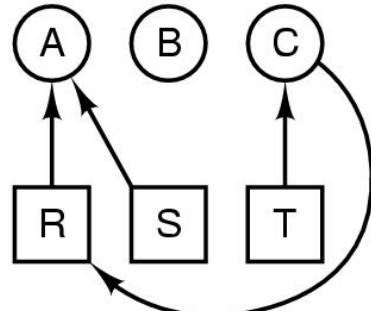
(l)



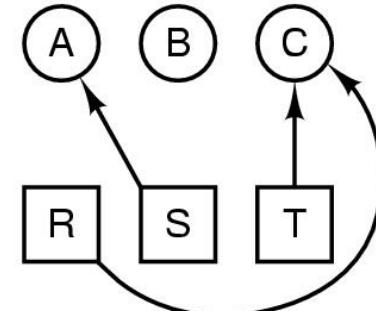
(m)



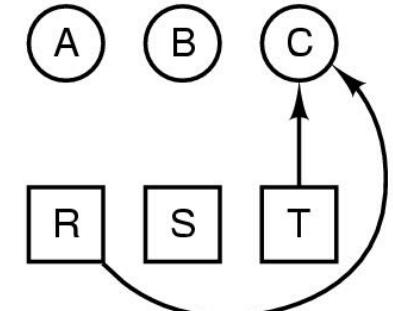
(n)



(o)



(p)



(q)

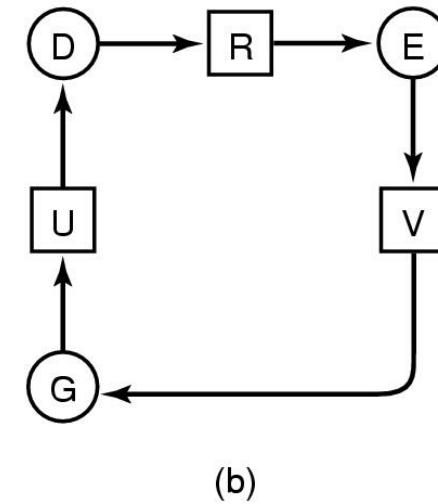
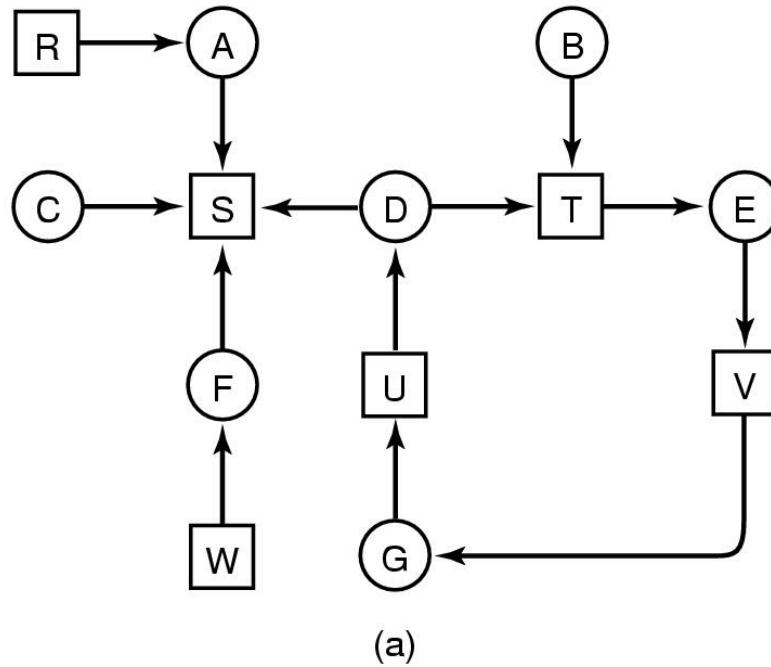
How deadlock can be avoided



The Ostrich Algorithm

- Pretend there is no problem
- Reasonable if
 - deadlocks occur very rarely
 - cost of prevention is high
- UNIX and Windows takes this approach
- It is a trade off between
 - convenience
 - correctness

Detection with One Resource of Each Type (1)



- Note the resource ownership and requests
- A cycle can be found within the graph, denoting deadlock

Detection with One Resource of Each Type (2)

Resources in existence
($E_1, E_2, E_3, \dots, E_m$)

Current allocation matrix

C_{11}	C_{12}	C_{13}	\dots	C_{1m}
C_{21}	C_{22}	C_{23}	\dots	C_{2m}
:	:	:		:
C_{n1}	C_{n2}	C_{n3}	\dots	C_{nm}

Row n is current allocation
to process n

Resources available
($A_1, A_2, A_3, \dots, A_m$)

Request matrix

R_{11}	R_{12}	R_{13}	\dots	R_{1m}
R_{21}	R_{22}	R_{23}	\dots	R_{2m}
:	:	:		:
R_{n1}	R_{n2}	R_{n3}	\dots	R_{nm}

Row 2 is what process 2 needs

Data structures needed by deadlock detection algorithm

Detection with One Resource of Each Type (3)

Tape drives
Plotters
Scanners
CD Roms

$$E = \begin{pmatrix} 4 & 2 & 3 & 1 \end{pmatrix}$$

Resources in existence

Tape drives
Plotters
Scanners
CD Roms

$$A = \begin{pmatrix} 2 & 1 & 0 & 0 \end{pmatrix}$$

Resources available

Current allocation matrix

$$C = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 1 & 2 & 0 \end{bmatrix}$$

Request matrix

$$R = \begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 2 & 1 & 0 & 0 \end{bmatrix}$$

An example for the deadlock detection algorithm

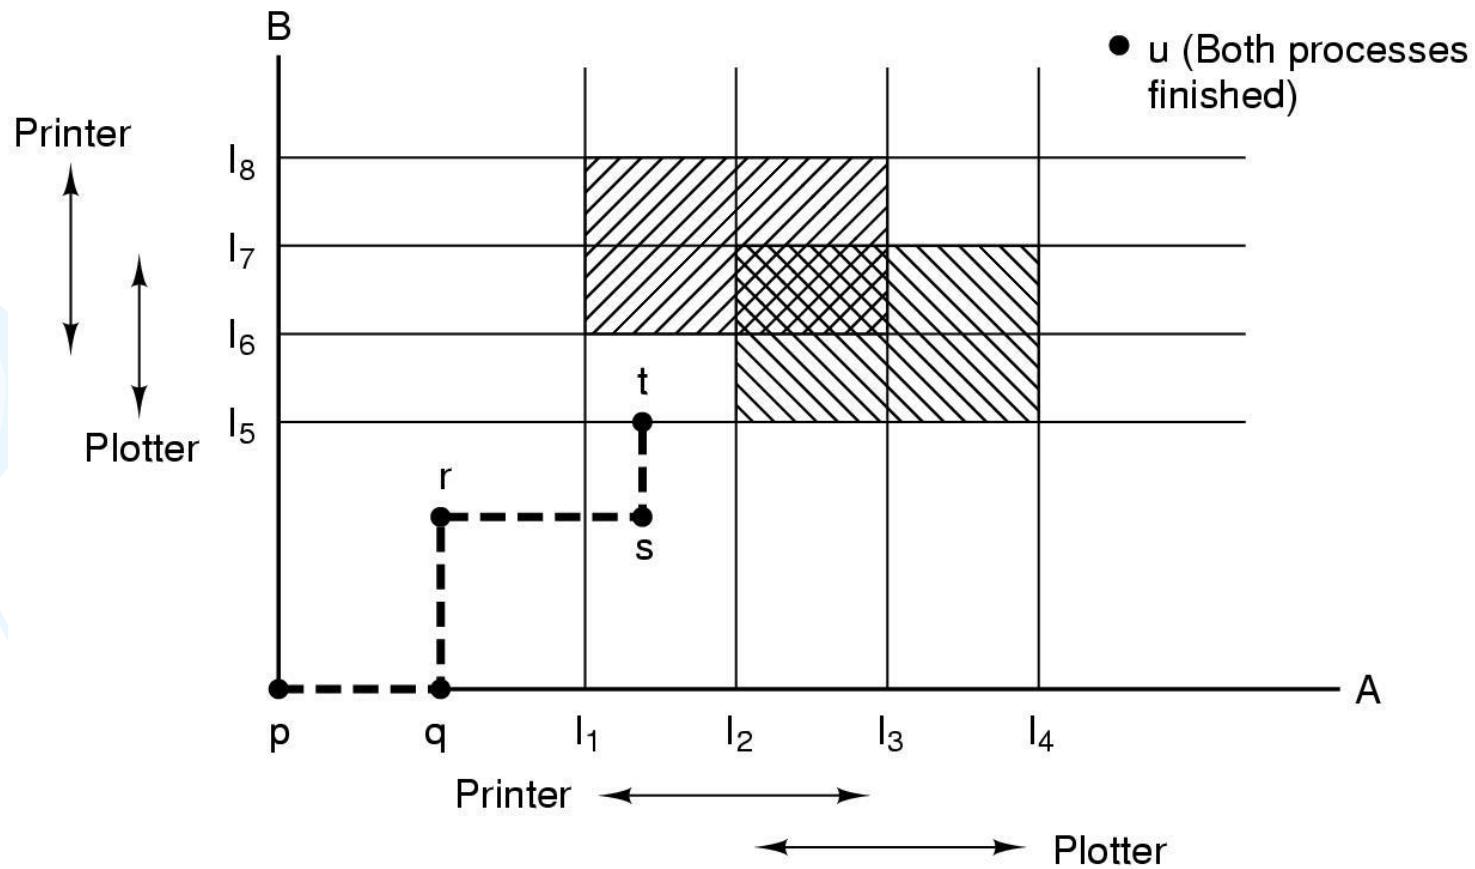
Recovery from Deadlock (1)

- Recovery through preemption
 - take a resource from some other process
 - depends on nature of the resource
- Recovery through rollback
 - checkpoint a process periodically
 - use this saved state
 - restart the process if it is found deadlocked

Recovery from Deadlock (2)

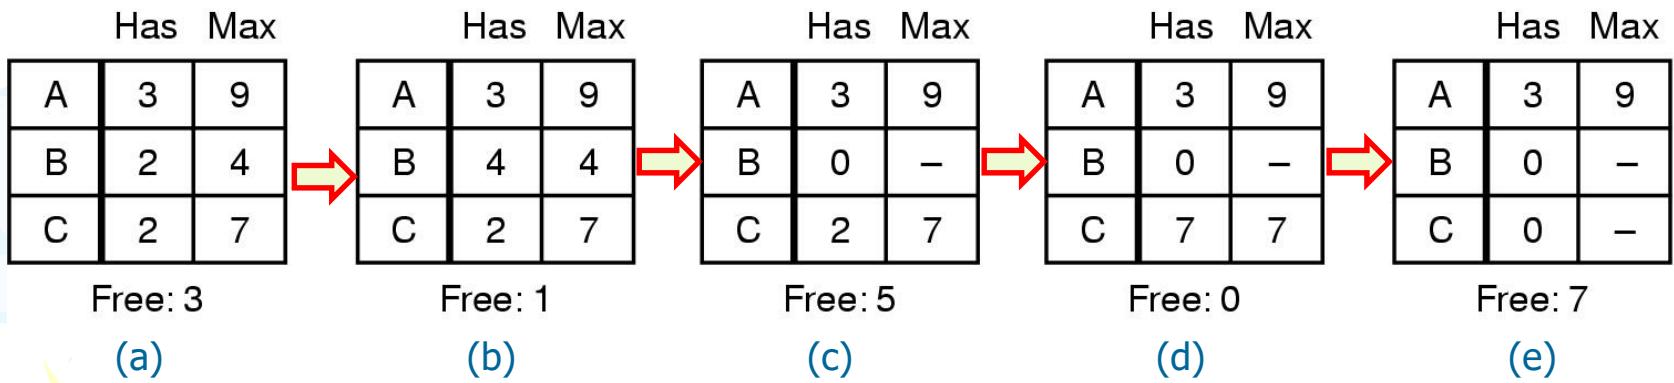
- Recovery through killing processes
 - crudest but simplest way to break a deadlock
 - kill one of the processes in the deadlock cycle
 - the other processes get its resources
 - choose process that can be rerun from the beginning

Deadlock Avoidance Resource Trajectories



Two process resource trajectories

Safe and Unsafe States (1)



Demonstration that the state in (a) is safe

Safe and Unsafe States (2)

	Has	Max
A	3	9
B	2	4
C	2	7

Free: 3

(a)

	Has	Max
A	4	9
B	2	4
C	2	7

Free: 2

(b)

	Has	Max
A	4	9
B	4	4
C	2	7

Free: 0

(c)

	Has	Max
A	4	9
B	—	—
C	2	7

Free: 4

(d)

Demonstration that the state in b is not safe

The Banker's Algorithm for a Single Resource (Dijkstra 1965)

	Has	Max
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

(a)

	Has	Max
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

(b)

	Has	Max
A	1	6
B	2	5
C	2	4
D	4	7

Free: 1

(c)

- Three resource allocation states
 - safe
 - safe
 - unsafe

Banker's Algorithm for Multiple Resources

Process
Tape drives
Plotters
Scanners
CD ROMs

A	3	0	1	1
B	0	1	0	0
C	1	1	1	0
D	1	1	0	1
E	0	0	0	0

Resources assigned

Process
Tape drives
Plotters
Scanners
CD ROMs

A	1	1	0	0
B	0	1	1	2
C	3	1	0	0
D	0	0	1	0
E	2	1	1	0

Resources still needed

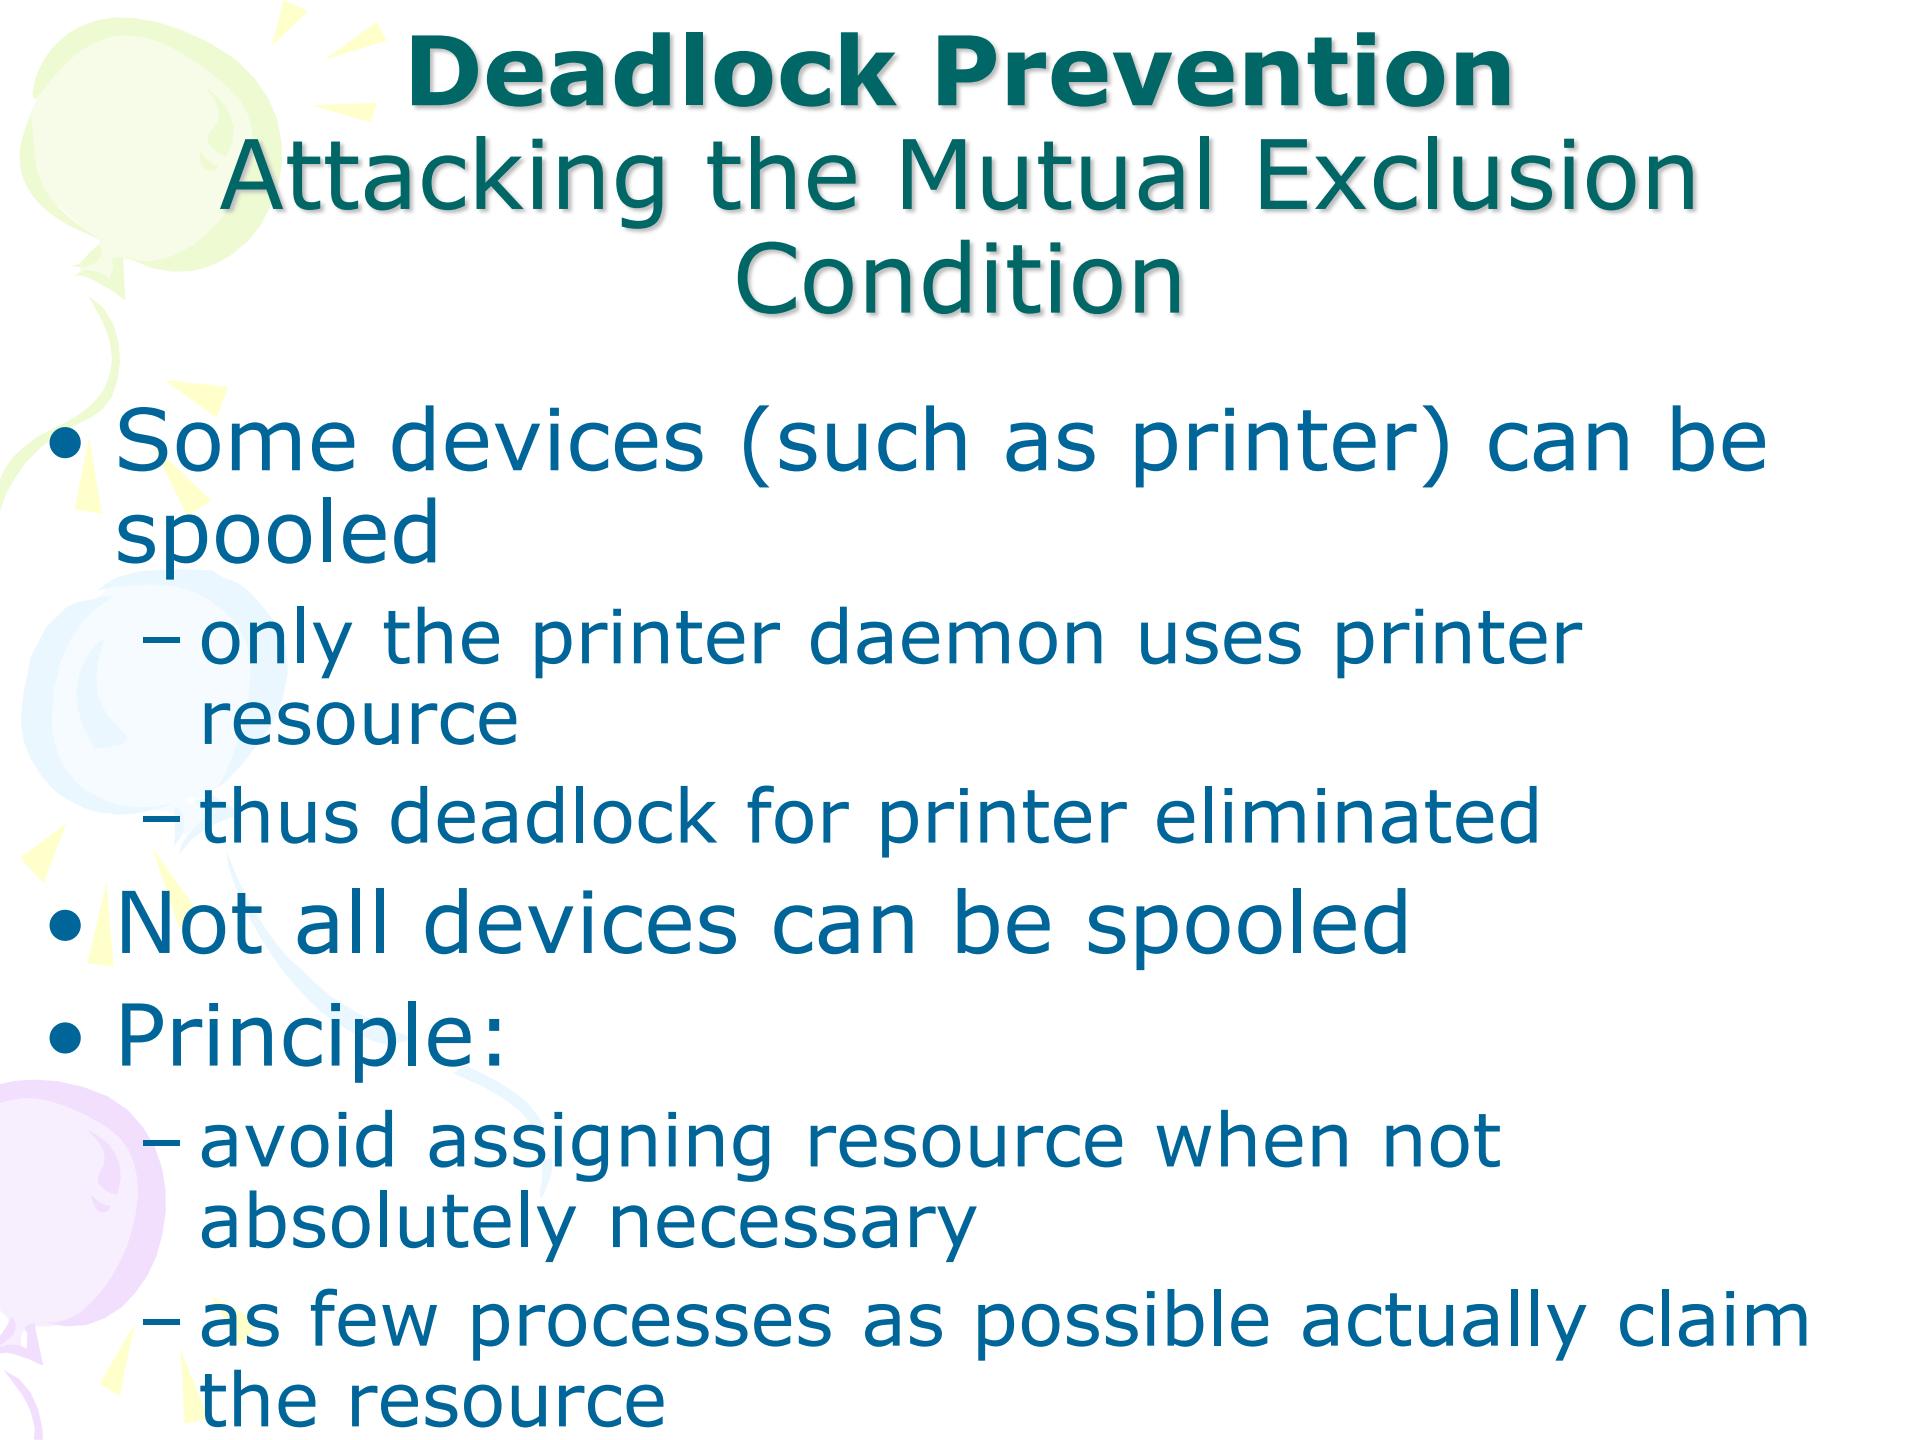
$E = (6342)$ Existing
 $P = (5322)$ Possessed=Resources
 $A = (1020)$ Available Assigned

Example of banker's algorithm with multiple resources



The Ostrich Algorithm

- Pretend there is no problem
- Reasonable if
 - deadlocks occur very rarely
 - cost of prevention is high
- UNIX and Windows takes this approach
- It is a trade off between
 - convenience
 - correctness



Deadlock Prevention

Attacking the Mutual Exclusion Condition

- Some devices (such as printer) can be spooled
 - only the printer daemon uses printer resource
 - thus deadlock for printer eliminated
- Not all devices can be spooled
- Principle:
 - avoid assigning resource when not absolutely necessary
 - as few processes as possible actually claim the resource

Deadlock Prevention

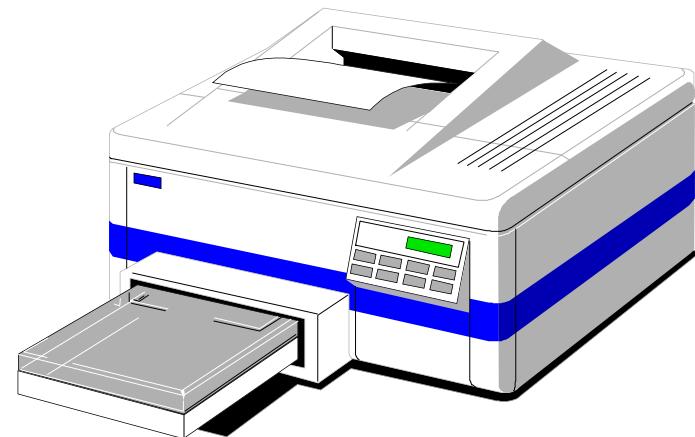
Attacking the Hold and Wait Condition

- Require processes to request resources before starting
 - a process never has to wait for what it needs
- Problems
 - may not know required resources at start of run
 - also ties up resources other processes could be using
- Variation:
 - process must give up all resources
 - then request all immediately needed

Deadlock Prevention

Attacking the No Preemption Condition

- This is not a viable option
- Consider a process given the printer
 - halfway through its job
 - now forcibly take away printer
 - !!??

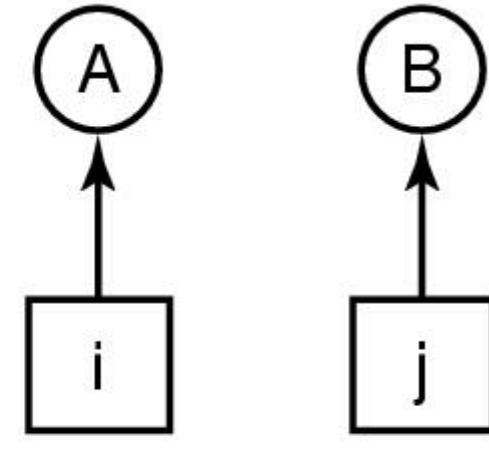


Deadlock Prevention

Attacking the Circular Wait Condition (1)

- 1. Imagesetter
- 2. Scanner
- 3. Plotter
- 4. Tape drive
- 5. CD Rom drive

(a)



(b)

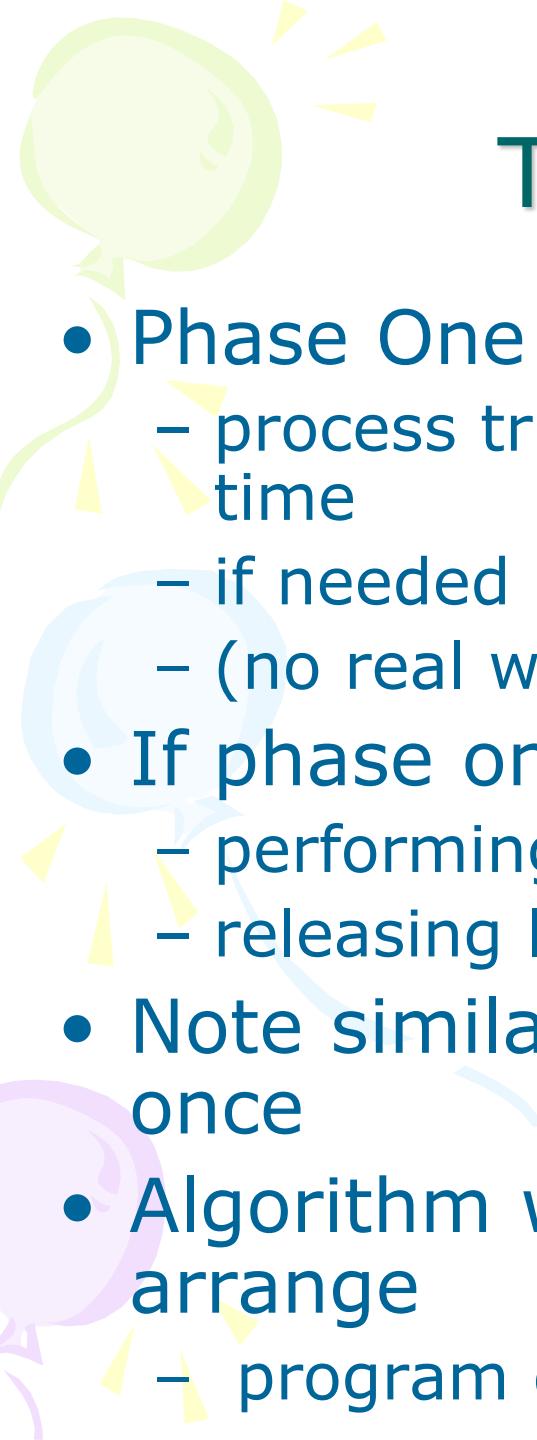
- a) Normally ordered resources
- b) A resource graph

Deadlock Prevention

Attacking the Circular Wait Condition (2)

Condition	Approach
Mutual exclusion	Spool everything
Hold and wait	Request all resources initially
No preemption	Take resources away
Circular wait	Order resources numerically

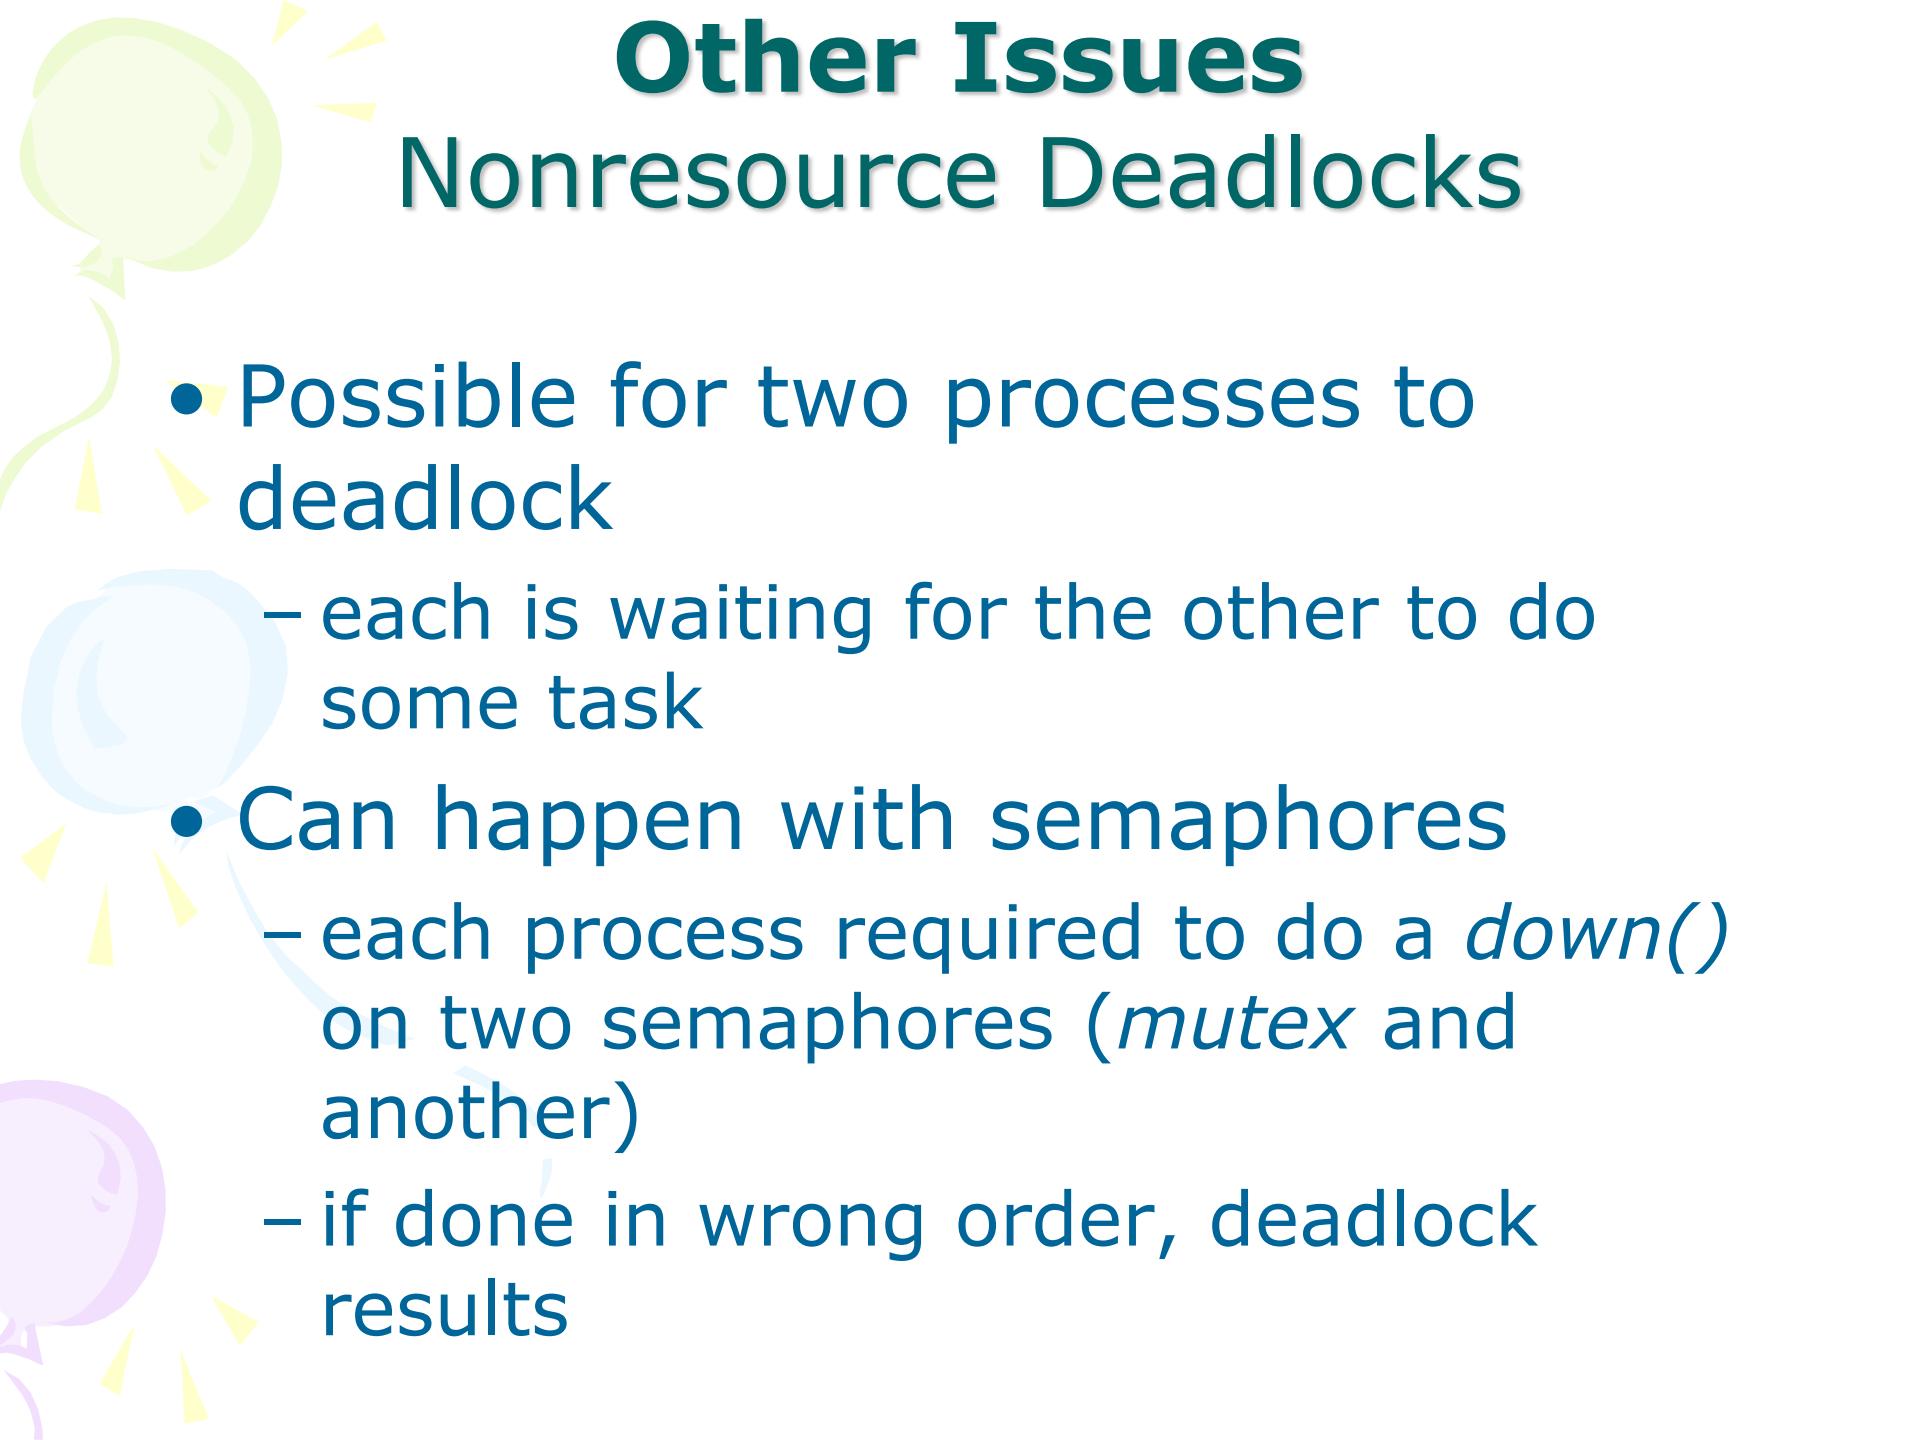
Summary of approaches to deadlock prevention



Other Issues

Two-Phase Locking

- Phase One
 - process tries to lock all records it needs, one at a time
 - if needed record found locked, start over
 - (no real work done in phase one)
- If phase one succeeds, it starts second phase,
 - performing updates
 - releasing locks
- Note similarity to requesting all resources at once
- Algorithm works where programmer can arrange
 - program can be stopped, restarted



Other Issues

Nonresource Deadlocks

- Possible for two processes to deadlock
 - each is waiting for the other to do some task
- Can happen with semaphores
 - each process required to do a *down()* on two semaphores (*mutex* and another)
 - if done in wrong order, deadlock results